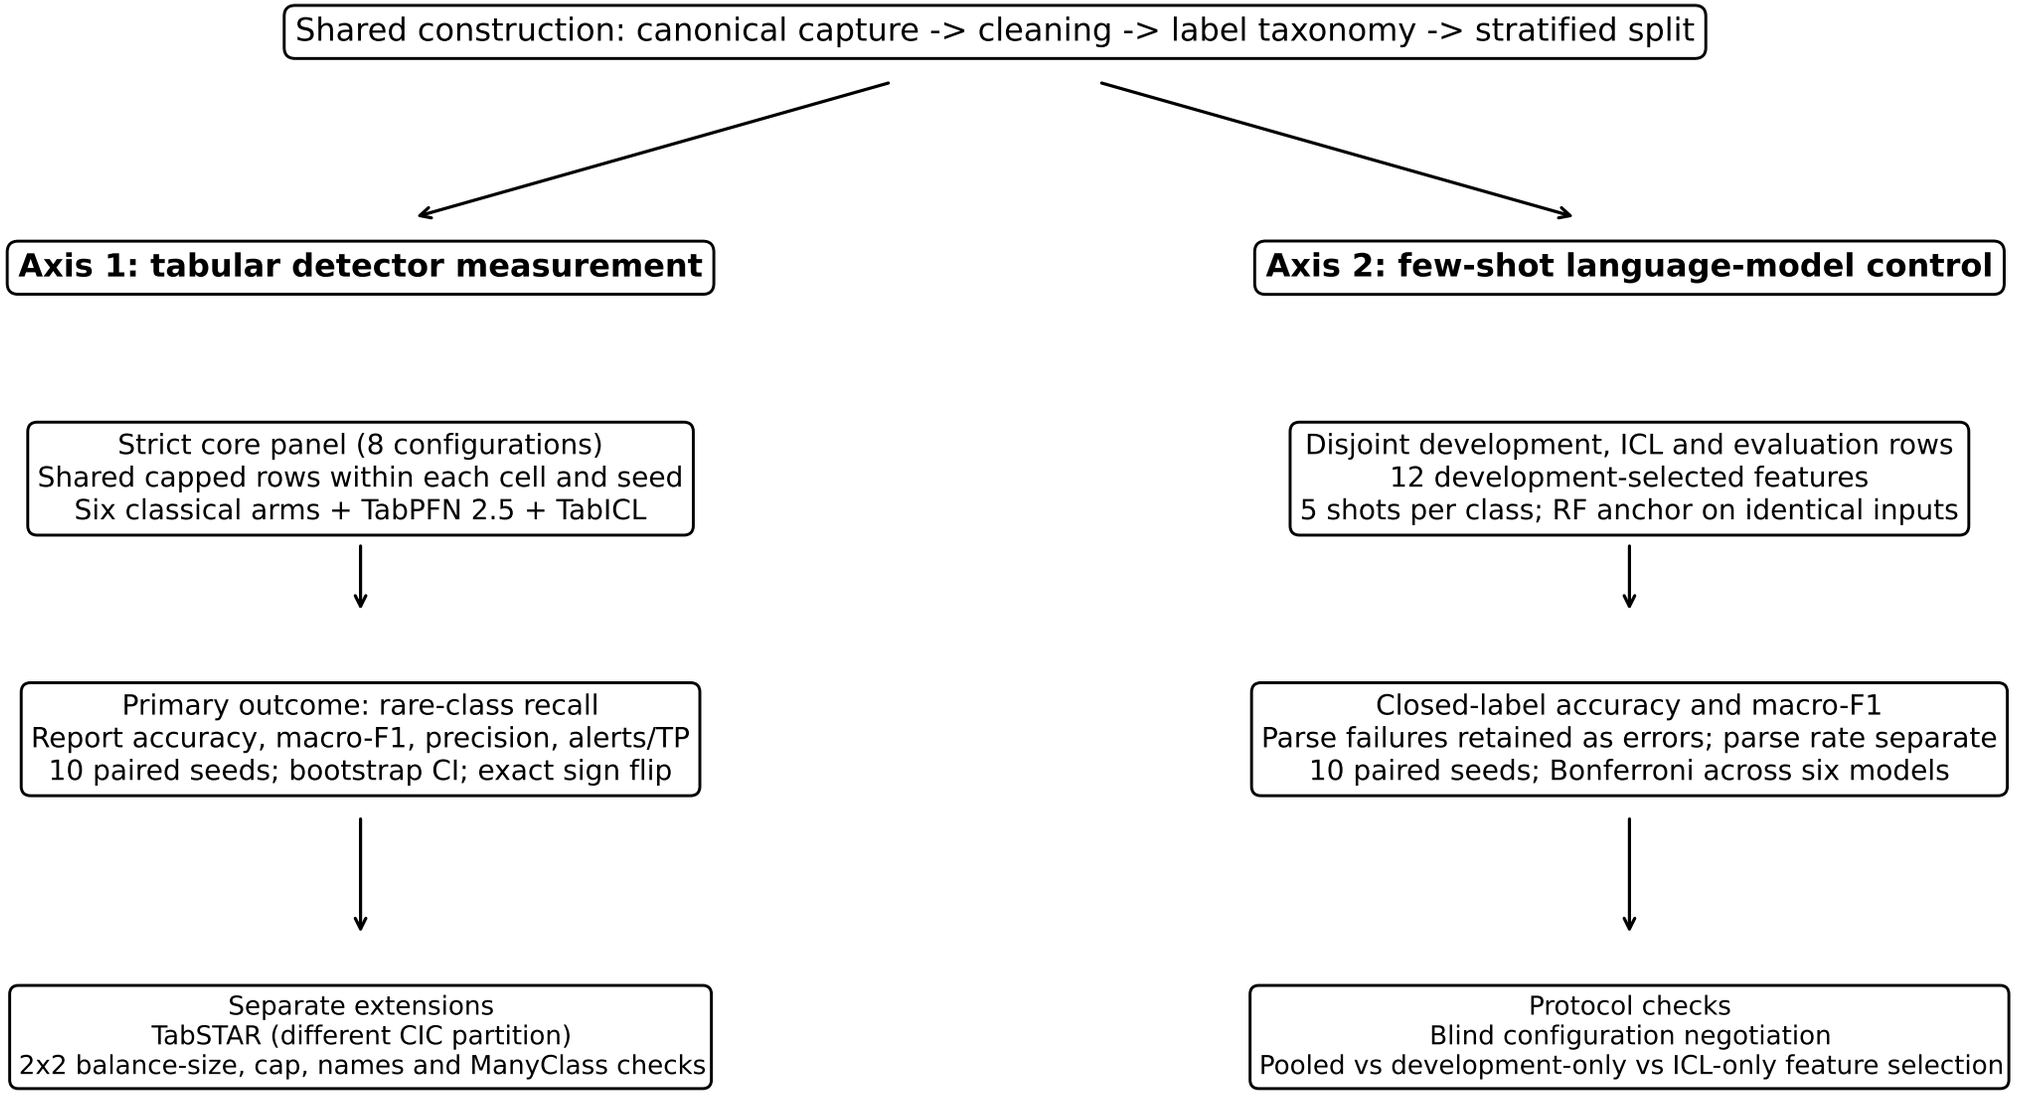
Shared construction: canonical capture -> cleaning -> label taxonomy -> stratified split
Axis 1: tabular detector measurement
Axis 2: few-shot language-model control
Strict core panel (8 configurations)
Shared capped rows within each cell and seed
Six classical arms + TabPFN 2.5 + TabICL
Disjoint development, ICL and evaluation rows
12 development-selected features
5 shots per class; RF anchor on identical inputs
Primary outcome: rare-class recall
Report accuracy, macro-F1, precision, alerts/TP
10 paired seeds; bootstrap CI; exact sign flip
Closed-label accuracy and macro-F1
Parse failures retained as errors; parse rate separate
10 paired seeds; Bonferroni across six models
Separate extensions
TabSTAR (different CIC partition)
2x2 balance-size, cap, names and ManyClass checks
Protocol checks
Blind configuration negotiation
Pooled vs development-only vs ICL-only feature selection