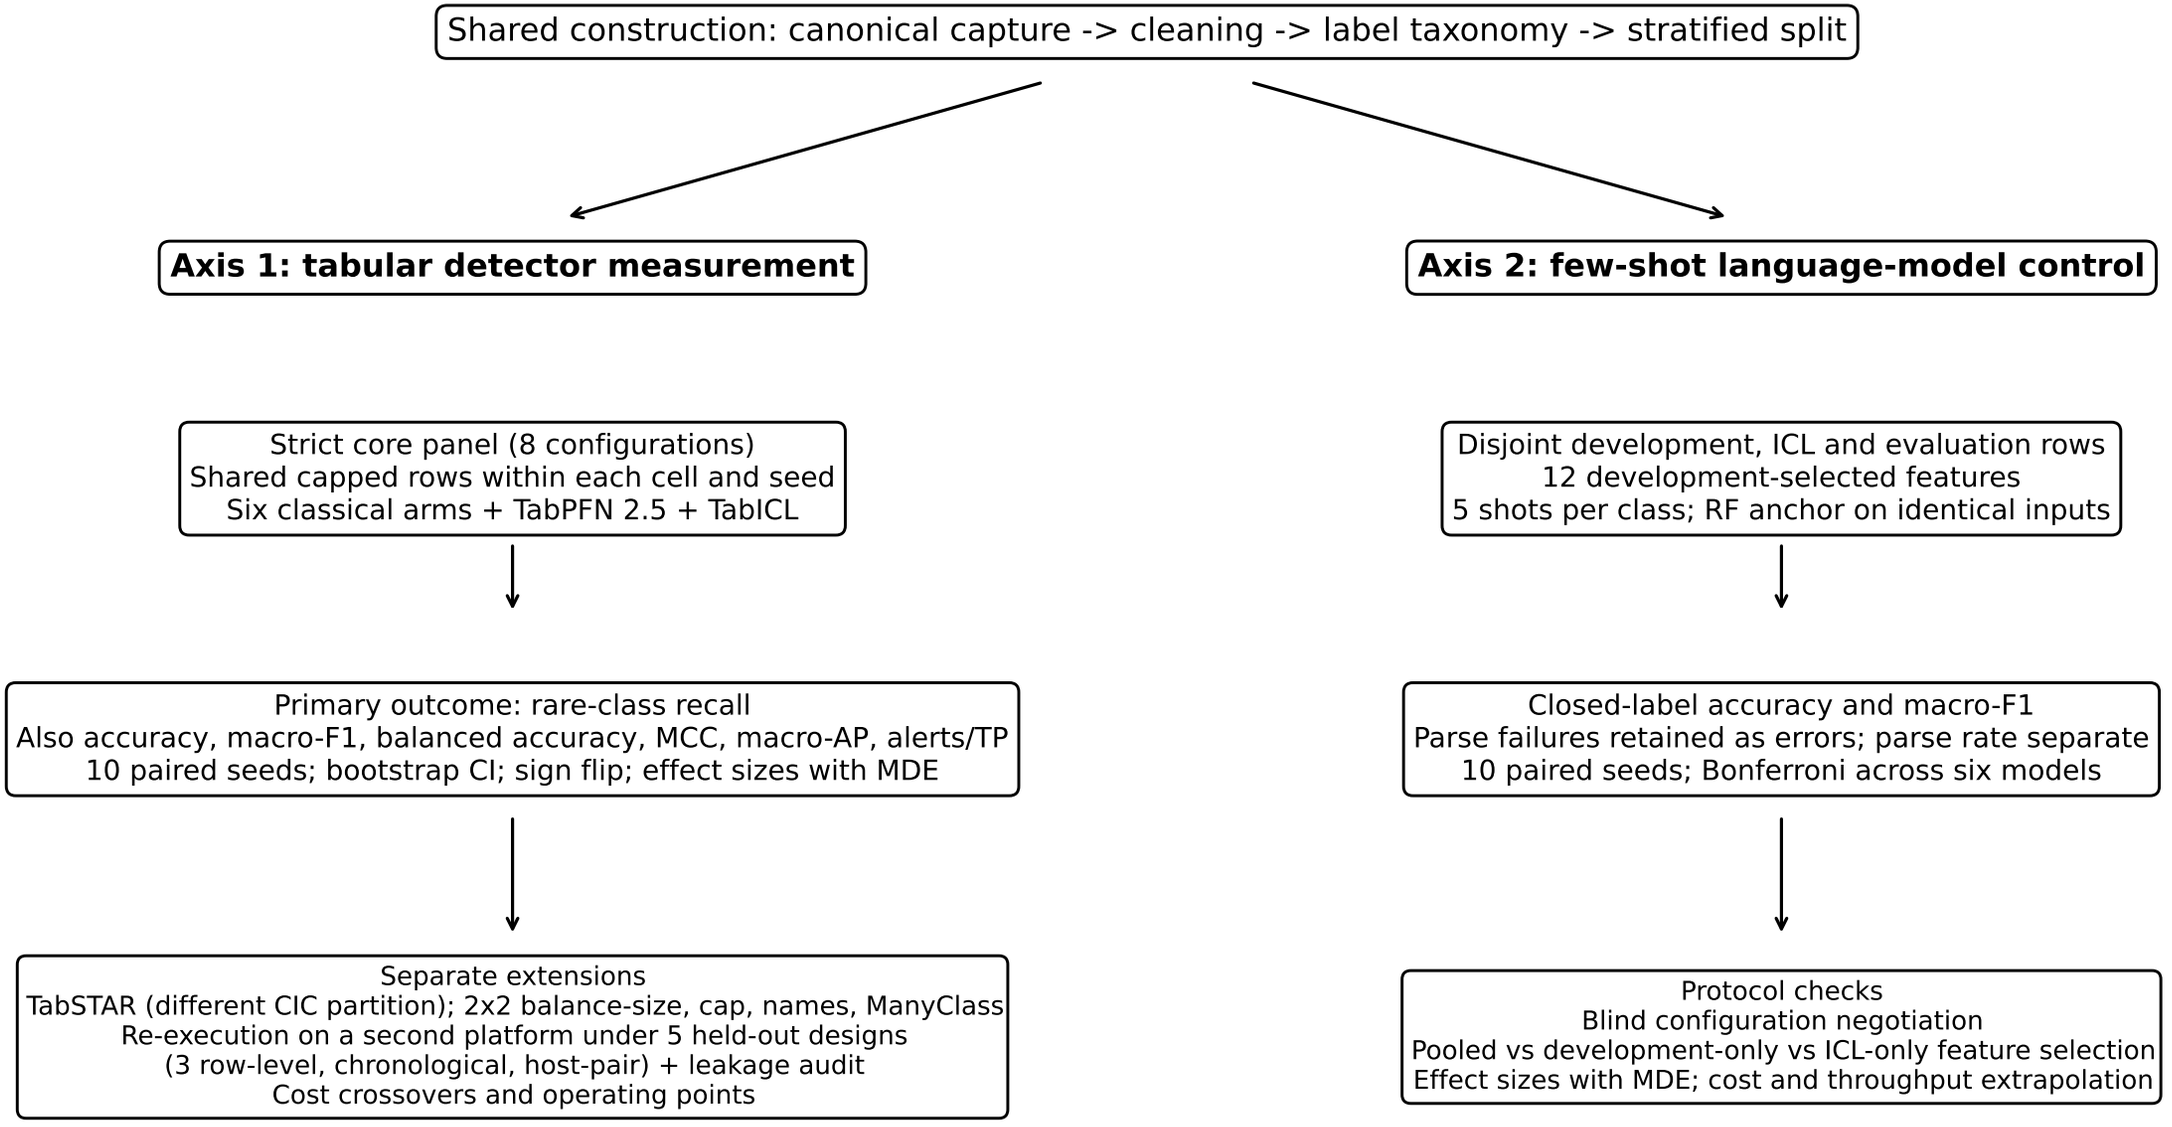
Shared construction: canonical capture -> cleaning -> label taxonomy -> stratified split
Axis 1: tabular detector measurement
Axis 2: few-shot language-model control
Strict core panel (8 configurations)
Shared capped rows within each cell and seed
Six classical arms + TabPFN 2.5 + TabICL
Disjoint development, ICL and evaluation rows
12 development-selected features
5 shots per class; RF anchor on identical inputs
Primary outcome: rare-class recall
Also accuracy, macro-F1, balanced accuracy, MCC, macro-AP, alerts/TP
10 paired seeds; bootstrap CI; sign flip; effect sizes with MDE
Closed-label accuracy and macro-F1
Parse failures retained as errors; parse rate separate
10 paired seeds; Bonferroni across six models
Separate extensions
TabSTAR (different CIC partition); 2x2 balance-size, cap, names, ManyClass
Re-execution on a second platform under 5 held-out designs
(3 row-level, chronological, host-pair) + leakage audit
Cost crossovers and operating points
Protocol checks
Blind configuration negotiation
Pooled vs development-only vs ICL-only feature selection
Effect sizes with MDE; cost and throughput extrapolation